

Erlang debugging

mats cronqvist

Ruminations on tools and strategies.

With boring anecdotes!

rumination

- n 1: a calm lengthy intent consideration [syn: contemplation, reflection, reflexion, musing, thoughtfulness]
- 2: (of ruminants) chewing (the cud); "ruminants have remarkable powers of rumination"
- 3: regurgitation of small amounts of food; seen in some infants after feeding

Masse goes to Denmark...

- Embedded system running OTP R5.
- Live in Denmark.
- There was no way to log into the CPU.
- There was no way to load new code.
- There was no usable application debugging tool.
- You could physically connect a terminal.

The node got overloaded after 90 days.

A tech traveled there and rebooted every 90 days.

...virtually...

- Wrote a one-liner...
- ...that ran a one-minute trace and wrote to a file.
- Sent it to the Danish tech by mail...
- ...who ran it by pasting it into a shell...
- ...before and after the reboot...
- ...and emailed the files to me (base-64 encoded)

...tracing saves the day!

Wrote a comparison profiler.

Compare the average execution time for each function, before and after the reboot.

`ets:lookup/2`

was 100 times slower before the reboot.

42

The hash function was broken for bignums.

The Dark Side

Embedded system.

Multi-million lines of C++.

The disks were too small for core files.

100s of units deployed.

...but...

The network worked.

The Point...

Debuggability is a property of a system

In a distributed system, fail-stop bugs are easy

The 3 kinds of bugs

- It crashes “randomly”
- It uses too much of a resource
- It gives the wrong answer

Strategy

	Monitoring	Narrowing
crashes	logging	context
performance	real-time logging	process → function
wrong result	contracts (test cases)	context

polling

The Erlang VM has many info methods

- `erlang:memory`
- `erlang:system_info`
- `erlang:statistics`
- `erlang:process_info`
- `inet:i()`.

Narrowing, procs

UNIX top

```
top - 13:54:15 up 24 days,  2:59,  9 users,  load average: 0.15, 0.42, 0.49
Tasks: 192 total,   5 running, 185 sleeping,   0 stopped,   2 zombie
Cpu(s):  7.0%us,  3.3%sy,  0.0%ni, 89.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem:   3106248k total,  2978996k used,   127252k free,    67844k buffers
Swap:          0k total,          0k used,          0k free,  2066100k cached
```

PID	%MEM	VIRT	SWAP	RES	CODE	DATA	SHR	nFLT	nDRT	S	PR	NI	%CPU	COMMAND
8748	4.5	181m	44m	136m	1544	174m	2156	15	0	S	20	0	1	beam.smp
2842	3.6	307m	196m	110m	44	236m	27m	0	0	R	20	0	7	firefox-bin
29380	3.0	395m	304m	91m	48	225m	30m	1095	0	S	20	0	0	amarok
24748	2.6	171m	93m	77m	31m	113m	12m	7	0	S	20	0	0	chrome

Narrowing, procs

dtop

dtop:start()

```
-----  
kr@sterlett          size: 4(45)M, cpu%: 0(27), procs: 39, runq: 0, 13:57:19  
memory[kB]:  proc      591, atom      300, bin      40, code    1905, ets      128
```

pid	name	current	msgq	mem	cpu
<0.49.0>	prfTarg	prfPrc:pidinfo/2	0	31	0
<0.40.0>	group:server/3	group:get_line1/3	0	11	0
<0.46.0>	dtop	prfHost:loop/1	0	21	0
<0.28.0>	user_drv:server/2	user_drv:server_loo	0	17	0
<0.50.0>	erlang:apply/2	file_io_server:serv	0	8	0
<0.51.0>	erlang:apply/2	file_io_server:serv	0	8	0

Interrupts

The Erlang VM has 2 interrupt mechanisms

- `erlang:trace/3` (redbug)
- `erlang:system_monitor/2` (watchdog)

Narrowing, functions

UNIX strace

STRACE(1)

STRACE(1)

NAME

strace - trace system calls and signals

SYNOPSIS

```
strace [ -dffhiqrstTvxx ] [ -acolumn ] [ -eexpr ] ... [ -ofile ] [
-ppid ] ... [ -sstrsize ] [ -uusername ] [ -Evar=val ] ... [ -Evar ]
... [ command [ arg ... ] ]
```

```
strace -c [ -eexpr ] ... [ -Ooverhead ] [ -Ssortby ] [ command [ arg ...
] ]
```

DESCRIPTION

Strace intercepts and records the system calls which are called by a process and the signals which are received by a process. The name of each system call, its arguments and its return value are printed...

Narrowing, functions

redbug

```
redbug:start(Trc,Opts).
```

redbug is a tool to interact with the Erlang trace facility. It will instruct the Erlang VM to generate so called 'trace messages' when certain events (such as a particular function being called) occur.

The trace messages are either printed (i.e. human readable) to a file or to the screen; or written to a trc file.

redbug trace patterns

```
Trc: list('send' | 'receive' | {M} | {M,F} | {M,RMSs} | {M,F,RMSs})
Proc: 'all' | pid() | atom(Regname) | {'pid', I2, I3}
Targ: node()
RMSs: (restricted match specs): list(RMS)
RMS: 'stack' | 'return' | tuple(ArgDescriptor)
ArgDescriptor: '_' | literal()
```

redbug opts

Opts: list({Opt,Val})

general opts:

time	(15000)	stop trace after this many ms
msgs	(10)	stop trace after this many msgs
proc	(all)	(list of) Erlang process(es)
targ	(node())	node to trace on

print-related opts

max_queue	(5000)	fail if internal queue gets this long
max_msg_size	(50000)	fail if seeing a msg this big
buffered	(no)	buffer messages till end of trace
print_form	("~s~n")	print msgs using this format
print_file	(standard_io)	print to this file
print_re	("")	print only strings that match this

trc file related opts

file	(none)	use a trc file based on this name
file_size	(1)	size of each trc file
file_count	(8)	number of trc files

redbug examples

```
> redbug:start({erlang},{msgs,1}).  
14:51:05 <kdb_sync> {erlang,'++',  
                    ["/home/masse/svn/kred/devel-xmlrepc/logs",  
                     "/TICK"]}
```

```
> redbug:start({erlang,'++'},[{msgs,1}]).  
14:51:48 <kdb_sync> {erlang,'++',  
                    ["/home/masse/svn/kred/devel-xmlrepc/logs",  
                     "/TICK"]}
```

```
> redbug:start({erlang,'++',[{"",'_'}]},{msgs,2}).  
14:52:54 <ssl_broker_sup> {erlang,'++',[[]],[<0.550.0>,acceptor,[[]]]}
```

more examples

```
> redbug:start({erlang,'++',[{"",'_'}]},{msgs,20},{print_file,"/tmp/foo"}).  
  
> redbug:start({erlang,'++'},[{msgs,2},{proc,kdb_sync}]).  
14:57:00 <kdb_sync> {erlang,'++',  
                    ["/home/masse/svn/kred/devel-xmlrepc/logs",  
                     "/TICK"]}
```

redbug for wizards

```
> redbug:start({erlang,'++',[stack,return]},{msgs,2},{proc,kdb_sync}).
14:58:36 <kdb_sync> {erlang,'++',
                    ["/home/masse/svn/kred/devel-xmlrepc/logs",
                     "/TICK"]}

{proc_lib,init_p_do_apply,3}
{gen_server,handle_msg,5}
{kdb_sync,handle_info,2}

14:58:36 <kdb_sync> {erlang,'++',2} -> "/home/masse/svn/kred/devel-
xmlrepc/logs/TICK"
```